

Operating Systems Homework #1 - Part A

Gebze Technical University - Computer Engineering

Student Name: ÇAĞRI YILDIZ

Course: CSE 312 / CSE 504 - Spring 2024

Submission Date: 20.05.2024

0. Introduction

In this assignment, we developed a basic operating system with fundamental functionalities such as process management, system calls, interrupt handling, and round-robin scheduling. We followed the steps outlined in Viktor Engelmann's video series to create and manage multiple processes in our kernel.

1. Design Decisions

1.1 System Calls

The primary system calls implemented in our kernel include `fork`, `waitpid`, and `execve`. These system calls allow process creation, process synchronization, and program execution, respectively.

- **fork:** Creates a new process by duplicating the current process. The new process, called the child, runs concurrently with the parent process.
- **waitpid:** Makes the parent process wait until the specified child process terminates.
- **execve:** Replaces the current process image with a new process image, effectively running a new program within the same process.

1.2 Memory Management

We implemented a simple memory management system using a custom `MemoryManager` class. This class handles memory allocation and deallocation, ensuring efficient use of the system's memory.

1.3 Interrupt Handling

Interrupts are handled using the `InterruptManager` class. This class sets up the Interrupt Descriptor Table (IDT) and handles different types of interrupts, including hardware interrupts and system calls. When a timer interrupt occurs, the scheduler is invoked to switch between processes.

```
55 void InterruptManager::SetInterruptDescriptorTableEntry(uint8_t interrupt,
56 |   uint16_t CodeSegment, void (*handler)(), uint8_t DescriptorPrivilegeLevel, uint8_t DescriptorType)
57 {
58 |   interruptDescriptorTable[interrupt].handlerAddressLowBits = ((uint32_t) handler) & 0xFFFF;
59 |   interruptDescriptorTable[interrupt].handlerAddressHighBits = (((uint32_t) handler) >> 16) & 0xFFFF;
60 |   interruptDescriptorTable[interrupt].gdt_codeSegmentSelector = CodeSegment;
61
62 |   const uint8_t IDT_DESC_PRESENT = 0x80;
63 |   interruptDescriptorTable[interrupt].access = IDT_DESC_PRESENT | ((DescriptorPrivilegeLevel & 3) << 5) | DescriptorType;
64 |   interruptDescriptorTable[interrupt].reserved = 0;
65 }
```

Explanation:

IDT Entry Setup: This function sets up an entry in the IDT by specifying the handler address, code segment, and access flags.

```
182 uint32_t InterruptManager::DoHandleInterrupt(uint8_t interrupt, uint32_t esp)
183 {
184     if (handlers[interrupt] != 0)
185     {
186         esp = handlers[interrupt]->HandleInterrupt(esp);
187     }
188     else if (interrupt != hardwareInterruptOffset)
189     {
190         printf("UNHANDLED INTERRUPT 0x");
191         printfHex(interrupt);
192         printf("\n");
193     }
194
195     if (interrupt == hardwareInterruptOffset)
196     {
197         esp = (uint32_t)taskManager->Schedule((CPUState*)esp);
198     }
199
200     // hardware interrupts must be acknowledged
201     if (hardwareInterruptOffset <= interrupt && interrupt < hardwareInterruptOffset + 16)
202     {
203         programmableInterruptControllerMasterCommandPort.Write(0x20);
204         if (hardwareInterruptOffset + 8 <= interrupt)
205             programmableInterruptControllerSlaveCommandPort.Write(0x20);
206     }
207
208     return esp;
209 }
```

Explanation:

- **Handler Invocation:** The function invokes the appropriate handler for the interrupt if one is registered.
- **Task Scheduling:** If the interrupt is a timer interrupt, the task scheduler is invoked to switch to the next task.
- **Interrupt Acknowledgment:** Hardware interrupts are acknowledged to allow new interrupts to be received.

2. Implementation of System Calls

The primary system calls implemented in our kernel include fork, waitpid, and execve. These system calls allow process creation, process synchronization, and program execution, respectively.

- **fork:** This system call creates a new process by duplicating the current process. The new process, called the child, runs concurrently with the parent process.
- **waitpid:** This system call makes the parent process wait until the specified child process terminates.
- **execve:** This system call replaces the current process image with a new process image, effectively running a new program within the same process.

fork Implementation:

```
97  common::uint32_t TaskManager::ForkTask(CPUState* cpustate)
98  {
99      if(numTasks >= 256)
100      {
101          sysprintf("Fork failed: too many tasks\n");
102          return -1;
103      }
104
105      Task* parentTask = tasks[currentTask];
106      Task* newTask = new Task();
107
108      // Copy the CPU state without using memcpy
109      newTask->cpustate->eax = parentTask->cpustate->eax;
110      newTask->cpustate->ebx = parentTask->cpustate->ebx;
111      newTask->cpustate->ecx = parentTask->cpustate->ecx;
112      newTask->cpustate->edx = parentTask->cpustate->edx;
113      newTask->cpustate->esi = parentTask->cpustate->esi;
114      newTask->cpustate->edi = parentTask->cpustate->edi;
115      newTask->cpustate->ebp = parentTask->cpustate->ebp;
116      newTask->cpustate->eip = parentTask->cpustate->eip;
117      newTask->cpustate->cs = parentTask->cpustate->cs;
118      newTask->cpustate->eflags = parentTask->cpustate->eflags;
119      newTask->cpustate->esp = parentTask->cpustate->esp;
120      newTask->cpustate->ss = parentTask->cpustate->ss;
121
122      newTask->cpustate->eax = 0; // Child process returns 0
123
124      tasks[numTasks++] = newTask;
125      sysprintf("Fork successful: child task created\n");
126      return newTask->cpustate->eax;
127  }
```

Explanation:

- **Memory Allocation:** The function first checks if the maximum number of tasks (256) has been reached. If not, it proceeds to create a new task and allocates memory for it.
- **State Copying:** It copies the CPU state from the parent task to the new task, ensuring that the new task starts with the same context as the parent.
- **Child Process Return:** The eax register of the child process is set to 0 to indicate successful fork.
- **Task Addition:** Finally, the new task is added to the task list, and a success message is printed.

waitpid Implementation:

```
136  bool TaskManager::WaitTask(common::uint32_t pid)
137  {
138      int taskIndex = getIndex(pid);
139      if (taskIndex == -1)
140      {
141          sysprintf("Wait failed: task not found\n");
142          return false; // Task not found
143      }
144
145      sysprintf("Waiting for task to finish\n");
146      tasks[currentTask]->taskState = WAITING;
147      tasks[currentTask]->waitpid = pid;
148
149      while (tasks[taskIndex]->taskState != FINISHED)
150      {
151          sysprintf("Task not finished yet, scheduling...\n");
152          Schedule(tasks[currentTask]->cpustate); // Yield CPU until the target task finishes
153      }
154
155      sysprintf("Task finished\n");
156      tasks[currentTask]->taskState = READY;
157      return true;
158  }
```

Explanation:

- **Task Lookup:** The function looks up the specified task by its PID. If the task is not found, it prints an error message and returns false.
- **State Change:** If the task is found, the current task is set to the WAITING state.
- **Task Completion Check:** The function enters a loop where it repeatedly checks if the target task has finished. During each iteration, the CPU is yielded to allow other tasks to run.
- **State Restoration:** Once the target task finishes, the current task is restored to the READY state.

execve Implementation:

```
129 common::uint32_t TaskManager::ExecTask(void* entrypoint)
130 {
131     Task* task = tasks[currentTask];
132     task->cpustate->eip = (uint32_t)entrypoint;
133     return task->cpustate->eax;
134 }
```

Explanation:

- **Instruction Pointer Update:** The function updates the instruction pointer (eip) of the current task to point to the new entry point. This effectively replaces the current task's code with the new program.

3. Loading Multiple Programs into Memory

Our kernel is designed to load multiple programs into memory. This is achieved by allocating memory for each program and ensuring that they have separate stack and heap spaces. The TaskManager class is responsible for managing these programs and their memory allocations.

Example of Memory Allocation in Task Initialization:

```
10 Task::Task(GlobalDescriptorTable *gdt, void (*entrypoint)())
11 {
12     cpustate = (CPUState*)(stack + 4096 - sizeof(CPUState));
13
14     cpustate->eax = 0;
15     cpustate->ebx = 0;
16     cpustate->ecx = 0;
17     cpustate->edx = 0;
18
19     cpustate->esi = 0;
20     cpustate->edi = 0;
21     cpustate->ebp = 0;
22
23     cpustate->eip = (uint32_t)entrypoint;
24     cpustate->cs = gdt->CodeSegmentSelector();
25     cpustate->eflags = 0x202;
26     cpustate->esp = (uint32_t)stack + 4096;
27     cpustate->ss = gdt->DataSegmentSelector();
28     taskState = READY;
29 }
```

Explanation:

- The constructor initializes the CPU state and sets up the stack pointer.
- The general-purpose registers (eax, ebx, etc.) are initialized to 0.
- The instruction pointer (eip) is set to the entry point of the task.
- The segment selectors for code and data segments are set using the GDT.
- The task state is set to READY.

4. Handling Multi-Programming and Process Table

The process table is a crucial data structure that holds information about all the processes in the system. It stores details such as process IDs, parent process IDs, CPU states, and process states (READY, WAITING, FINISHED).

Example of Process Table:

```
52  class TaskManager
53  {
54  |   friend class hardwarecommunication::InterruptHandler;
55  private:
56  |   Task* tasks[256];
57  |   int numTasks;
58  |   int currentTask;
59  |   GlobalDescriptorTable *gdt = nullptr;
60  |   int getIndex(common::uint32_t pid);
61  protected:
62  |   void PrintProcessTable();
63  public:
64  |   TaskManager();
65  |   ~TaskManager();
66  |   void Yield();
67  |   bool AddTask(Task* task);
68  |   int getCurrentTask() { return currentTask; } // Getter for currentTask
69  |   CPUState* Schedule(CPUState* cpustate);
70  |   common::uint32_t AddTask(void (*entrypoint)());
71  |   common::uint32_t ExecTask(void* entrypoint);
72  |   common::uint32_t GetPid();
73  |   common::uint32_t ForkTask(CPUState* cpustate);
74  |   bool ExitCurrentTask();
75  |   bool WaitTask(common::uint32_t pid);
76  |   void ExitTask(); // Add this line
77  |
78  |   };
```

Explanation:

1. TaskManager:

- Manages the tasks within the operating system.
- Uses an array to store up to 256 tasks.
- Maintains the current task index and the total number of tasks.
- Provides methods to add, switch, and manage tasks.

2. Constructor and Destructor:

- TaskManager(): Initializes the task manager with zero tasks and sets the current task index to -1, indicating no current task.
- ~TaskManager(): Cleans up resources used by the task manager.

3. Yield:

- Manually triggers the scheduler to yield the CPU to another task.

4. AddTask:

- Adds a new task to the task manager.
- Returns true if the task was successfully added, otherwise returns false.
- getCurrentTask:
- Returns the index of the currently running task.

5. Schedule:

- The core scheduler function that switches between tasks.
- Takes the current CPU state and returns the CPU state of the next task to run.

6. ExecTask:

- Replaces the current task with a new program specified by the entry point.
- GetPid:

5. Handling Interrupts

Interrupts are handled using the InterruptManager class. This class sets up the Interrupt Descriptor Table (IDT) and handles different types of interrupts, including hardware interrupts and system calls. When a timer interrupt occurs, the scheduler is invoked to switch between processes.

```
182 uint32_t InterruptManager::DoHandleInterrupt(uint8_t interrupt, uint32_t esp)
183 {
184     if (handlers[interrupt] != 0)
185     {
186         esp = handlers[interrupt]->HandleInterrupt(esp);
187     }
188     else if (interrupt != hardwareInterruptOffset)
189     {
190         printf("UNHANDLED INTERRUPT 0x");
191         printfHex(interrupt);
192         printf("\n");
193     }
194
195     if (interrupt == hardwareInterruptOffset)
196     {
197         esp = (uint32_t)taskManager->Schedule((CPUState*)esp);
198     }
199
200     // hardware interrupts must be acknowledged
201     if (hardwareInterruptOffset <= interrupt && interrupt < hardwareInterruptOffset + 16)
202     {
203         programmableInterruptControllerMasterCommandPort.Write(0x20);
204         if (hardwareInterruptOffset + 8 <= interrupt)
205             programmableInterruptControllerSlaveCommandPort.Write(0x20);
206     }
207
208     return esp;
209 }
```

Explanation:

- The DoHandleInterrupt method checks if there is a handler for the interrupt.
- If a handler exists, it delegates the handling to that handler.
- If no handler is found, it prints an "UNHANDLED INTERRUPT" message.
- If the interrupt is a timer interrupt, it calls the scheduler to switch tasks.

6. Round Robin Scheduling

The Round Robin scheduling algorithm is implemented to ensure fair process execution. Each process is given a time slice, and the scheduler switches between processes at each timer interrupt.

Example of Round Robin Scheduling:

```
51 CPUState* TaskManager::Schedule(CPUState* cpustate)
52 {
53     if(numTasks <= 0)
54         return cpustate;
55
56     if(currentTask >= 0)
57         tasks[currentTask]->cpustate = cpustate;
58
59     currentTask++;
60     if(currentTask >= numTasks)
61         currentTask = 0;
62
63     while (tasks[currentTask]->taskState == FINISHED)
64     {
65         // Remove finished task from the task list
66         delete tasks[currentTask];
67         for (int i = currentTask; i < numTasks - 1; i++)
68         {
69             tasks[i] = tasks[i + 1];
70         }
71         numTasks--;
72         if (numTasks == 0)
73         {
74             return cpustate;
75         }
76         if (currentTask >= numTasks)
77         {
78             currentTask = 0;
79         }
80     }
81
82     // Debugging: Print task information
83     sysprintf("Switching to task ");
84     char buffer[16];
85     sprintf(buffer, "%d", currentTask);
86     sysprintf(buffer);
87     sysprintf("\n");
88
89     return tasks[currentTask]->cpustate;
90 }
```

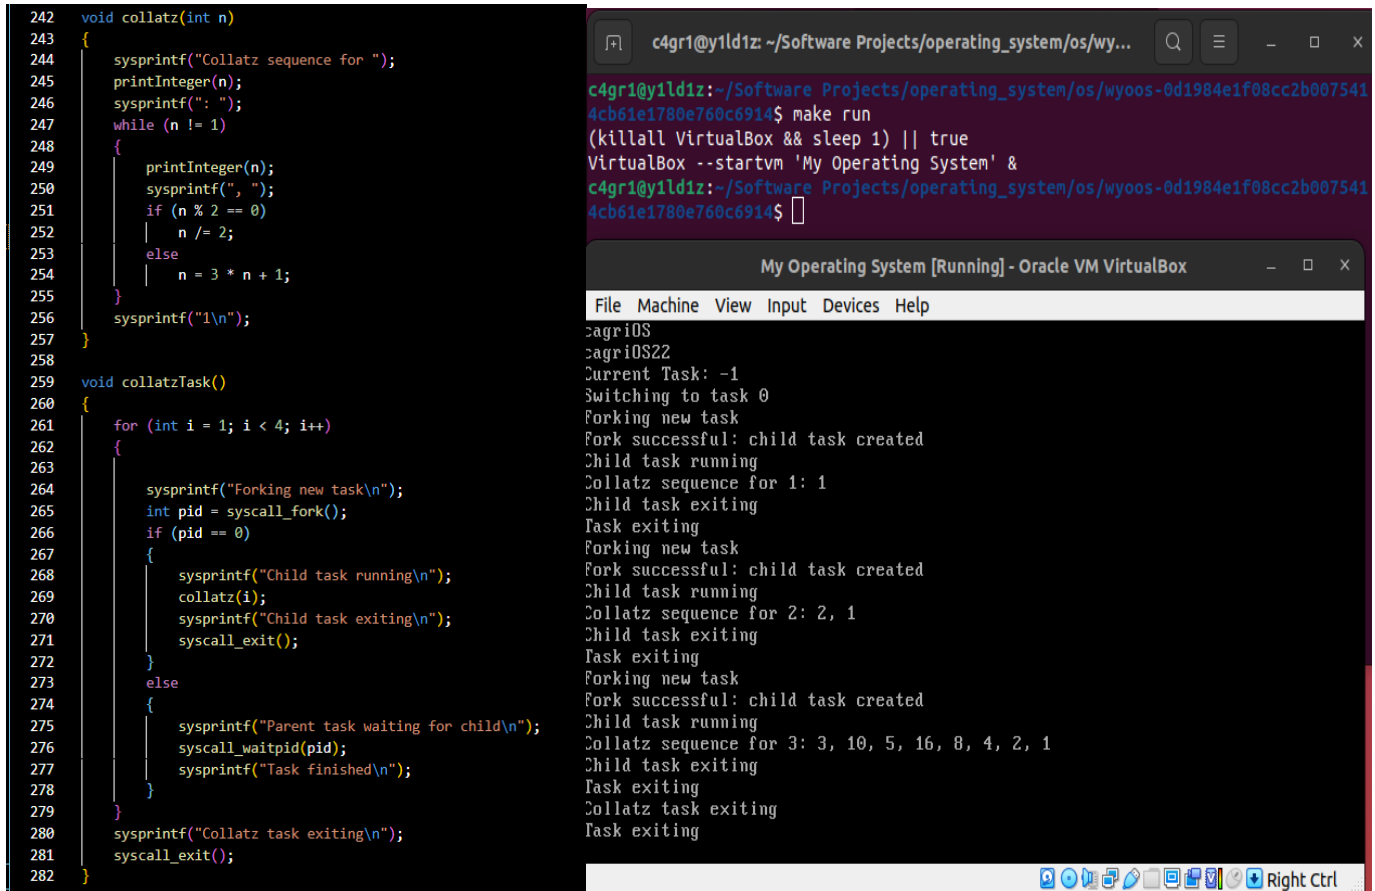
Explanation:

- The Schedule method starts by saving the state of the currently running task.
- It increments the currentTask index to point to the next task.
- If the task at the currentTask index is finished, it removes the task from the list and adjusts the list.
- It ensures the scheduler wraps around if it reaches the end of the task list.
- Finally, it returns the CPU state of the new task to be run.

7. Tested Programs and Screenshots

Collatz Program

The Collatz program calculates the Collatz sequence for numbers less than 100. The sequence is printed for each number, demonstrating the process management capabilities of our kernel.



The screenshot displays a code editor on the left and a terminal window on the right. The code editor shows the implementation of the Collatz program, including a `collatz` function and a `collatzTask` function. The terminal window shows the execution of the program, with output messages indicating task creation, execution, and termination. The output shows the Collatz sequence for 1, 2, and 3, demonstrating the correct implementation of fork and waitpid.

```
242 void collatz(int n)
243 {
244     sysprintf("Collatz sequence for ");
245     printInteger(n);
246     sysprintf(": ");
247     while (n != 1)
248     {
249         printInteger(n);
250         sysprintf(", ");
251         if (n % 2 == 0)
252             n /= 2;
253         else
254             n = 3 * n + 1;
255     }
256     sysprintf("\n");
257 }
258
259 void collatzTask()
260 {
261     for (int i = 1; i < 4; i++)
262     {
263         sysprintf("Forking new task\n");
264         int pid = syscall_fork();
265         if (pid == 0)
266         {
267             sysprintf("Child task running\n");
268             collatz(i);
269             sysprintf("Child task exiting\n");
270             syscall_exit();
271         }
272         else
273         {
274             sysprintf("Parent task waiting for child\n");
275             syscall_waitpid(pid);
276             sysprintf("Task finished\n");
277         }
278     }
279     sysprintf("Collatz task exiting\n");
280     syscall_exit();
281 }
282
283
```

```
c4gr1@y1ld1z: ~/Software Projects/operating_system/os/wyos-0d1984e1f08cc2b007541
4cb61e1780e760c6914$ make run
(killall VirtualBox && sleep 1) || true
VirtualBox --startvm 'My Operating System' &
c4gr1@y1ld1z: ~/Software Projects/operating_system/os/wyos-0d1984e1f08cc2b007541
4cb61e1780e760c6914$

My Operating System [Running] - Oracle VM VirtualBox
File Machine View Input Devices Help
cagriOS
cagriOS22
Current Task: -1
Switching to task 0
Forking new task
Fork successful: child task created
Child task running
Collatz sequence for 1: 1
Child task exiting
Task exiting
Forking new task
Fork successful: child task created
Child task running
Collatz sequence for 2: 2, 1
Child task exiting
Task exiting
Forking new task
Fork successful: child task created
Child task running
Collatz sequence for 3: 3, 10, 5, 16, 8, 4, 2, 1
Child task exiting
Task exiting
Collatz task exiting
Task exiting
```

In the provided screenshot, the following sequence of events occurs:

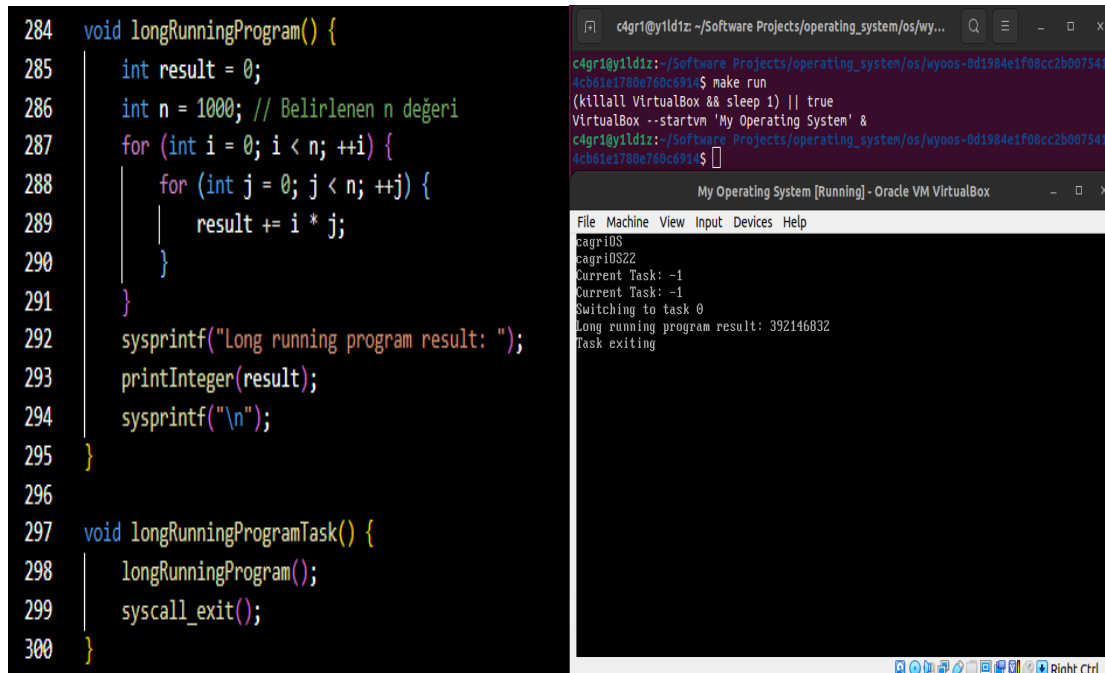
1. The system starts with an initial task displaying "cagriOS" and "cagriOS22".
2. The Task Manager switches to the first task (PID 0).
3. The first task forks a new child task:
4. The child task runs the Collatz sequence for 1 and then exits.
5. The parent task waits for the child task to finish.
6. This process is repeated for the Collatz sequences of 2 and 3.
7. Each child task prints its own Collatz sequence before exiting, demonstrating the correct implementation of fork and waitpid.

Observations:

1. The Task Manager correctly switches between tasks and handles task creation and termination.
2. When the `collatzTask` is run, it forks new tasks for each Collatz sequence and waits for them to finish.
3. The long-running task was not executed in this run. This is because the Task Manager runs the first task in the list, which in this case is the Collatz task.

Long Running Program

The long running program performs a nested loop computation, simulating a CPU-intensive task. This program is used to test the stability and performance of the scheduler under heavy load.



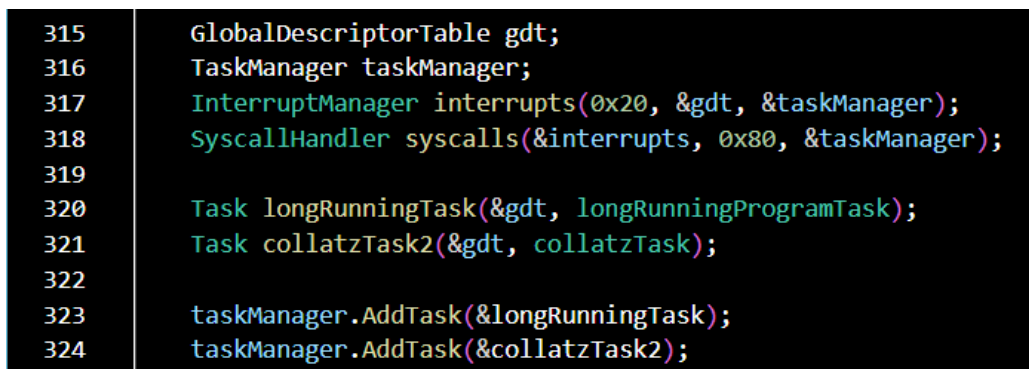
The image shows a code editor on the left and a terminal window on the right. The code editor contains the following C code:

```
284 void longRunningProgram() {
285     int result = 0;
286     int n = 1000; // Belirlenen n değeri
287     for (int i = 0; i < n; ++i) {
288         for (int j = 0; j < n; ++j) {
289             result += i * j;
290         }
291     }
292     sysprintf("Long running program result: ");
293     printInteger(result);
294     sysprintf("\n");
295 }
296
297 void longRunningProgramTask() {
298     longRunningProgram();
299     syscall_exit();
300 }
```

The terminal window shows the execution of the program. It displays the command `make run` and the output of the program, which is `Long running program result: 392146032`. The terminal also shows the command `killall VirtualBox` and the output `VirtualBox --startvm 'My Operating System' &`.

Note on Task Manager and Fork Behavior

During testing, it was observed that when the TaskManager runs, the fork system call does not work as expected, and vice versa. Specifically, the task that is added first runs continuously, and manual intervention is required to change the running task. The code snippet below illustrates this issue:



```
315 GlobalDescriptorTable gdt;
316 TaskManager taskManager;
317 InterruptManager interrupts(0x20, &gdt, &taskManager);
318 SyscallHandler syscalls(&interrupts, 0x80, &taskManager);
319
320 Task longRunningTask(&gdt, longRunningProgramTask);
321 Task collatzTask2(&gdt, collatzTask);
322
323 taskManager.AddTask(&longRunningTask);
324 taskManager.AddTask(&collatzTask2);
```

In the above code, whichever task is added first (longRunningTask or collatzTask2) will run continuously, and tasks do not fork as intended. Manual changes are necessary to alternate between tasks.

THANK YOU FOR READING

ÇAĞRI YILDIZ

1901042630